

MICRO PROJECT

**DEEP LEARNING
(M24CS1E203C)**

Bayesian Deep Learning on MNIST using EM-TDAMP

MICRO PROJECT REPORT

Submitted by

PAUL JOSE

MAC24CSCE07

To

*The APJ ABDUL KALAM TECHNOLOGICAL UNIVERSITY in partial fulfillment for the award
of the degree*

of

MASTER OF TECHNOLOGY

IN

COMPUTER SCIENCE AND ENGINEERING



DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

MAR ATHANASIOUS COLLEGE OF ENGINEERING

(GOVT. AIDED & AUTONOMOUS)

KOTHAMANGALAM, KERALA-686 666

APRIL 2025

DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING
MAR ATHANASIUS COLLEGE OF ENGINEERING
(GOVT. AIDED & AUTONOMOUS)
KOTHAMANGALAM, KERALA-686 666



CERTIFICATE

This is to certify that the report entitled “**Bayesian Deep Learning on MNIST using EM-TDAMP**” submitted by **Mr. Paul Jose , Reg No: MAC24CSCE07** to the APJ Abdul Kalam Technological University in partial fulfillment of the requirements for the award of the Degree of Master of Technology in Computer Science & Engineering for the academic year 2024-2025 is a bonafied record of the micro project presented by them under our supervision and guidance. This report in any form has not been submitted to any other university or Institute for any purpose.

.....

Prof. Pristy Paul T
Staff in Charge

.....

Prof. Joby George
Head of the Department

ACKNOWLEDGEMENT

First and foremost, I sincerely thank **God Almighty** for his grace for the successful and timely completion of the micro project. I express my sincere gratitude and thanks to the Principal **Dr. Bos Mathew Jos** and Head of the Department **Prof. Joby George** for providing the necessary facilities and their encouragement and support. I owe special thanks to the faculty in charge **Prof. Pristly Paul T** for their corrections, suggestions and efforts to coordinate the micro project under a tight schedule. I also express my gratitude to the staff members in the Department of Computer Science and Engineering who have taken sincere efforts in helping me to completing this microproject. Finally, I would like to acknowledge the tremendous support given to me by our dear friends without whose support this work would have been all the more difficult to accomplish.

ABSTRACT

This project investigates the application of the Expectation Maximization and Turbo Deep Approximate Message Passing (EM-TDAMP) algorithm for training a sparse Multi-Layer Perceptron (MLP) on the MNIST dataset. By reformulating deep neural network learning as a sparse Bayesian inference problem, EM-TDAMP achieves structured model compression while maintaining high classification performance. The implemented MLP, with a 784-128-10 architecture, is trained on 10,000 MNIST samples, targeting 50% neuron sparsity in the hidden layer. After 10 EM iterations, the model attains a validation accuracy of 92.77%, with 64 active neurons, demonstrating a balance between efficiency and performance. Metrics such as precision (0.9294), recall (0.9263), and F1 score (0.9268) indicate robust classification, while low predictive entropy (0.0070) reflects high confidence. Compared to standard MLPs, the sparse model performs competitively, highlighting EM-TDAMP's potential for resource-constrained applications. This work underscores the value of Bayesian deep learning in creating compact, effective neural networks.

Contents

List of Figures	i
1 Introduction	1
2 System Design	2
3 Program	4
4 Result	13
5 Conclusion	17
REFERENCES	18

List of Figures

4.1	Classification accuracy versus training iteration. The model steadily improves its predictive performance, with accuracy converging as training progresses.	14
4.2	Evolution of precision, recall, and F1-score over training iterations. These metrics highlight the balance between false positives and false negatives across the learning process.	15
4.3	Confusion matrix of the final trained model. The diagonal dominance indicates strong class-wise accuracy, while off-diagonal elements reveal misclassification trends.	15
4.4	Number of active neurons versus iteration during training. The plot illustrates how the model prunes redundant units over time, reflecting the influence of Bayesian regularization.	15
4.5	Distribution of the variational parameter ρ , representing posterior uncertainty in the model's weights. A wider spread indicates higher uncertainty, a hallmark of Bayesian deep learning.	16
4.6	Relationship between network sparsity and classification accuracy. The plot demonstrates that high sparsity can be achieved with minimal degradation in performance, emphasizing the model's efficiency.	16
4.7	Per-class uncertainty estimates. Classes with higher uncertainty may correspond to ambiguous inputs or data imbalance, guiding further analysis and potential dataset refinement.	16

CHAPTER 1

INTRODUCTION

Deep learning has transformed fields like image recognition and natural language processing, but it faces challenges such as overfitting, high computational costs, and large model sizes that hinder deployment on resource-constrained devices. Bayesian deep learning offers a promising solution by incorporating uncertainty quantification and sparsity, leading to more robust and efficient models. This approach models neural network parameters as probability distributions, enabling structured compression and improved generalization compared to traditional methods.

This project leverages the Expectation Maximization and Turbo Deep Approximate Message Passing (EM-TDAMP) algorithm, introduced in the paper "Bayesian Deep Learning via Expectation Maximization and Turbo Deep Approximate Message Passing" (arXiv). The EM-TDAMP algorithm reformulates deep neural network (DNN) learning and compression as a sparse Bayesian inference problem. By employing group sparse priors, it achieves structured model compression, pruning entire neurons or layers, and accelerates convergence compared to stochastic gradient descent (SGD)-based methods. The algorithm's applications extend to federated learning, but this project focuses on its use in a centralized setting for classification.

The EM-TDAMP algorithm is applied to the MNIST dataset, a standard benchmark for handwritten digit recognition. The MNIST database contains 60,000 training images and 10,000 testing images, each a 28x28 grayscale image of handwritten digits from 0 to 9 (MNIST Database). This dataset is widely used to evaluate machine learning algorithms due to its simplicity and well-understood characteristics. The project aims to train a sparse neural network that maintains high classification accuracy while reducing model complexity, demonstrating the practical utility of Bayesian methods.

CHAPTER 2

SYSTEM DESIGN

Model Architecture

The project implements a Multi-Layer Perceptron (MLP) designed for MNIST classification, with the following architecture:

- Input Layer: 784 neurons, corresponding to the flattened 28x28 pixel images.
- Hidden Layer: 128 neurons with ReLU (Rectified Linear Unit) activation, introducing non-linearity.
- Output Layer: 10 neurons, representing the 10 digit classes (0–9), with linear outputs processed by a probit-product likelihood function for classification.

The MLP is built using TensorFlow/Keras, allowing access to weights and activations necessary for the EM-TDAMP algorithm's message-passing operations.

Dataset and Preprocessing

The MNIST dataset is preprocessed to ensure compatibility with the model:

- Training Set: Limited to 10,000 samples (out of 60,000) for computational efficiency.
- Test Set: Uses the full 10,000 samples.
- Normalization: Pixel values are scaled to the range $[0, 1]$ by dividing by 255.
- Reshaping: Images are flattened from 28x28 2D arrays to 784-element 1D arrays.
- Label Encoding: Labels are one-hot encoded into 10-dimensional vectors for multi-class classification.

- **Data Pipeline:** Training data is shuffled, batched (batch size of 100), and prefetched using TensorFlow's `tf.data` API for efficient processing. The test data is batched similarly but not shuffled.

Training Process

The EM-TDAMP algorithm combines Expectation Maximization (EM) with Turbo Deep Approximate Message Passing (TDAMP) to train the MLP while enforcing sparsity. Key components include:

- **EM Framework:** The algorithm iterates between:
 - **E-step:** Estimating posterior distributions of model parameters using TDAMP.
 - **M-step:** Updating hyperparameters, such as noise variance and sparsity parameters.
- **TDAMP Modules:**
 - **Module A (DAMP):** Performs deep approximate message passing, sampling weights from Gaussian priors, computing logits, and updating posteriors using the Adam optimizer. The probit-product likelihood models classification uncertainty.
 - **Module B (SPMP):** Enforces group sparsity by updating sparsity parameters (ρ), pruning neurons based on their posterior means.
- **Training Parameters:**
 - **Batch Size:** 100 samples.
 - **EM Iterations:** 10 iterations.
 - **Target Sparsity:** 0.5, aiming to prune 50% of the hidden layer neurons (64 out of 128 remain active).
 - **Noise Variance:** Initialized at 1.0 and reduced by a factor of 0.99 per iteration (minimum $1e-6$) to enhance stability.
- **Sparsity Enforcement:** The algorithm thresholds sparsity parameters to achieve the target sparsity, setting smaller values to zero, effectively pruning neurons.
- **Evaluation:** After each EM iteration, the model is evaluated on the test set, computing metrics like accuracy, precision, recall, F1 score, and average predictive entropy. Early stopping is implemented with a patience of 3 iterations and a minimum accuracy improvement of 0.001.

The training process is logged extensively, capturing loss values, sparsity levels, and performance metrics, with visualizations generated for accuracy trends and confusion matrices.

CHAPTER 3

PROGRAM

```
import tensorflow as tf
from tensorflow import keras
import tensorflow_probability as tfp
import numpy as np
import scipy
import matplotlib.pyplot as plt
import argparse
import os
import logging
import pandas as pd
from sklearn.metrics import precision_recall_fscore_support, confusion_matrix
import seaborn as sns

# Set up logging
logging.basicConfig(
    level=logging.INFO,
    format='%(asctime)s - %(levelname)s - %(message)s'
)

logger = logging.getLogger(__name__)

with tf.device('/GPU:0'):
    a = tf.constant([1.0])
    b = tf.constant([2.0])
    c = a + b
    logger.info("GPU test result: %s", c.numpy())

# Set random seeds for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

# Parse command-line arguments
parser = argparse.ArgumentParser(description='EM-TDAMP for MNIST Classification')
parser.add_argument('--batch_size', type=int, default=100, help='Batch size for training')
parser.add_argument('--em_iters', type=int, default=50, help='Number of EM iterations')
parser.add_argument('--sparsity', type=float, default=0.5, help='Target sparsity level')
args = parser.parse_args()

# 1. Project Setup & Data
def load_and_preprocess_data():
```

```

"""Load and preprocess the full MNIST dataset."""
logger.info("Loading and preprocessing MNIST dataset")
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()

# Use full training data
x_train = x_train[:10000] / 255.0 # Take first 10,000 samples
y_train = y_train[:10000]
x_train = x_train.reshape(-1, 784)
x_test = x_test.reshape(-1, 784) / 255.0

# One-hot encode labels
y_train = keras.utils.to_categorical(y_train, 10)
y_test = keras.utils.to_categorical(y_test, 10)

train_dataset = tf.data.Dataset.from_tensor_slices(
    (x_train, y_train)).shuffle(60000).batch(args.batch_size).prefetch(tf.data.AUTOTUNE)
test_dataset = tf.data.Dataset.from_tensor_slices(
    (x_test, y_test)).batch(args.batch_size).prefetch(tf.data.AUTOTUNE)

logger.info("MNIST dataset loaded: %d training samples, %d test samples", x_train.shape[0], x_test.shape[0])
return train_dataset, test_dataset, x_train, y_train, x_test, y_test

# 2. Model Definition
def build_model():
    """Define MLP model with access to weights and activations."""
    logger.info("Building MLP model: 784 -> 128 (ReLU) -> 10")
    inputs = keras.Input(shape=(784,))
    hidden = keras.layers.Dense(128, activation='relu', name='hidden')(inputs)
    outputs = keras.layers.Dense(10, name='outputs')(hidden)
    model = keras.Model(inputs=inputs, outputs=outputs)
    return model

# 3. Probit-Product Likelihood
def probit_product_likelihood(logits, y_true, v):
    """Compute log probit-product likelihood for classification using vectorized operations."""
    y_idx = tf.argmax(y_true, axis=1) # Shape: [batch_size]
    y_idx = tf.cast(y_idx, tf.int32) # Ensure int32 dtype to match tf.range
    zi = logits # Shape: [batch_size, num_classes]
    yi = tf.gather(zi, y_idx, batch_dims=1) # Shape: [batch_size]
    diff = (zi - tf.expand_dims(yi, axis=1)) / tf.sqrt(tf.cast(v, tf.float32)) # Shape: [batch_size, num_classes]
    q_values = 0.5 * tf.math.erfc(diff / tf.sqrt(2.0)) # Shape: [batch_size, num_classes]
    indices = tf.stack([tf.range(tf.shape(logits)[0]), y_idx], axis=1) # Shape: [batch_size, 2]
    q_values = tf.tensor_scatter_nd_update(q_values, indices, tf.ones([tf.shape(logits)[0]], dtype=tf.float32))
    log_q = tf.math.log(q_values + 1e-10) # Shape: [batch_size, num_classes]
    log_likelihood = tf.reduce_sum(log_q, axis=1) # Shape: [batch_size]
    return log_likelihood

# 4. Turbo Deep Approximate Message-Passing (TDAMP)
class TDAMP:
    def __init__(self, model, num_groups, v_init=1.0):
        """Initialize TDAMP with model weights and Bayesian parameters."""
        logger.info("Initializing TDAMP with %d groups, initial v=%.2f", num_groups, v_init)

```

```

self.model = model

self.num_groups = num_groups # Number of neurons in hidden layer

self.v = tf.Variable(v_init, dtype=tf.float32, trainable=False) # Noise variance

self.rho = tf.Variable(tf.ones([num_groups], dtype=tf.float32) * 0.5, trainable=False) # Sparsity

parameters

# Initialize weight means and variances as Variables
self.weight_means = [tf.Variable(w, trainable=False) for w in model.get_weights()]
self.weight_variances = [tf.Variable(tf.ones_like(w) * 0.1, trainable=False) for w in model.get_weights()]
# Initialize legacy Adam optimizer for M1/M2 compatibility
self.optimizer = tf.keras.optimizers.legacy.Adam(learning_rate=0.01)

def module_a_damp(self, x, y, priors):
    """Module A: Deep Approximate Message Passing (simplified)."""
    # Unpack priors (means and variances)
    prior_means, prior_variances = priors

    # Sample weights from Gaussian priors
    sampled_weights = []
    for mean, var in zip(prior_means, prior_variances):
        noise = tf.random.normal(tf.shape(mean), mean=0.0, stddev=tf.sqrt(var))
        sampled_weight = mean + noise
        sampled_weights.append(sampled_weight)

    # Set model weights to sampled weights for forward pass
    original_weights = self.model.get_weights()
    self.model.set_weights(sampled_weights)

    # Compute logits and likelihood
    with tf.GradientTape() as tape:
        logits = self.model(x, training=False)
        log_likelihood = probit_product_likelihood(logits, y, self.v)
        loss = -tf.reduce_mean(log_likelihood)

    # Compute gradients w.r.t. trainable variables
    trainable_vars = self.model.trainable_variables
    gradients = tape.gradient(loss, trainable_vars)

    # Log gradient norms for debugging
    grad_norms = [tf.norm(g).numpy() if g is not None else 0.0 for g in gradients]
    logger.debug("Gradient norms: %s", grad_norms)

    # Update posterior means using Adam optimizer
    posterior_means = []
    posterior_variances = []
    for i, (mean, var, grad) in enumerate(zip(prior_means, prior_variances, gradients)):
        if grad is None:
            # Handle case where gradient is None (e.g., for biases if not used)
            posterior_means.append(mean)
            posterior_variances.append(var)
            continue

        # Create a variable for optimization
        temp_mean = tf.Variable(mean, trainable=True)

```

```

        self.optimizer.apply_gradients([(grad, temp_mean)])
        # Update variance: Reduce variance to simulate posterior tightening
        new_var = var * 0.95
        posterior_means.append(tf.Variable(temp_mean, trainable=False))
        posterior_variances.append(tf.Variable(new_var, trainable=False))

    # Restore original weights
    self.model.set_weights(original_weights)

    # Update internal weight means and variances
    for i in range(len(self.weight_means)):
        self.weight_means[i].assign(posterior_means[i])
        self.weight_variances[i].assign(posterior_variances[i])

    return posterior_means, posterior_variances, loss

def module_b_smp(self, posteriors):
    """Module B: Sum-Product Message Passing for group sparse priors (simplified)."""
    posterior_means, posterior_variances = posteriors

    # Simplified SPMP: Update sparsity (rho) based on posterior means
    hidden_weights_mean = posterior_means[0] # Shape: [784, 128]
    group_norms = tf.reduce_mean(tf.abs(hidden_weights_mean), axis=0) # Shape: [128]

    # Update rho: Higher group norm -> higher probability of being active
    new_rho = tf.sigmoid(group_norms * 10.0) # Scale for sharper probabilities
    self.rho.assign(tf.clip_by_value(new_rho, 0.0, 1.0))

    # Compute extrinsic messages (simplified: scale posteriors)
    extrinsic_means = [m * 0.99 for m in posterior_means]
    extrinsic_variances = [v * 0.99 for v in posterior_variances]

    return extrinsic_means, extrinsic_variances, self.rho

def step(self, x, y, batch_idx):
    """Perform one TDAMP iteration."""
    priors = (self.weight_means, self.weight_variances)
    posterior_means, posterior_variances, loss = self.module_a_damp(x, y, priors)
    posteriors = (posterior_means, posterior_variances)
    extrinsics, new_rho = self.module_b_smp(posteriors)[:2]
    self.weight_means = [tf.Variable(m, trainable=False) for m in posteriors[0]]
    self.weight_variances = posteriors[1]

    if batch_idx % 100 == 0:
        avg_rho = tf.reduce_mean(self.rho).numpy()
        logger.info("Batch %d: Loss = %.4f, Average rho = %.4f", batch_idx, loss, avg_rho)

    return posteriors

# 5. EM Loop
def em_loop(model, train_dataset, test_dataset, em_iters, sparsity_target):
    """Run EM-TDAMP training loop with early stopping."""

```

```

logger.info("Starting EM-TDAMP training for %d iterations with target sparsity %.2f", em_iters,
            sparsity_target)
tdamp = TDAMP(model, num_groups=128, v_init=1.0)
history = {
    'val_accuracy': [], 'val_precision': [], 'val_recall': [], 'val_f1': [],
    'sparsity': [], 'active_neurons': [], 'avg_entropy': []
}
patience = 3
min_delta = 0.001
best_accuracy = -float('inf')
patience_counter = 0

for em_iter in range(em_iters):
    logger.info("EM Iteration %d/%d", em_iter + 1, em_iters)
    batch_idx = 0
    total_loss = 0.0
    num_batches = 0

    for x_batch, y_batch in train_dataset:
        posteriors = tdamp.step(x_batch, y_batch, batch_idx)
        batch_idx += 1
        num_batches += 1
        total_loss += -tf.reduce_mean(probit_product_likelihood(model(x_batch, training=False), y_batch,
tdamp.v)).numpy()

    tdamp.v.assign(tf.maximum(tdamp.v * 0.99, 1e-6))
    threshold = np.percentile(tdamp.rho.numpy(), 100 * (1 - sparsity_target))
    threshold = tf.cast(threshold, tf.float32)
    tdamp.rho.assign(tf.where(tdamp.rho > threshold, tdamp.rho, 0.0))

    accuracy, precision, recall, f1, cm = evaluate_model(model, test_dataset, tdamp.weight_means)
    sparsity = compute_sparsity(tdamp.rho)
    avg_entropy = compute_average_entropy(model, test_dataset, tdamp.weight_means)
    avg_loss = total_loss / num_batches
    active_neurons = tf.reduce_sum(tf.cast(tdamp.rho > 0, tf.float32)).numpy()

    history['val_accuracy'].append(float(accuracy))
    history['val_precision'].append(float(precision))
    history['val_recall'].append(float(recall))
    history['val_f1'].append(float(f1))
    history['sparsity'].append(float(sparsity))
    history['active_neurons'].append(int(active_neurons))
    history['avg_entropy'].append(float(avg_entropy))

    if accuracy > best_accuracy + min_delta:
        best_accuracy = accuracy
        patience_counter = 0
        final_cm = cm # Save the confusion matrix of the best model
    else:
        patience_counter += 1
        if patience_counter >= patience:
            logger.info("Early stopping triggered at iteration %d", em_iter + 1)

```

```

        break

    logger.info(
        "Val Accuracy = %.4f, Precision = %.4f, Recall = %.4f, F1 = %.4f, Sparsity = %.4f, Avg Entropy = %.4f, Avg Loss = %.4f, Active Neurons = %d/%d",
        accuracy, precision, recall, f1, sparsity, avg_entropy, avg_loss, int(active_neurons),
        tdamp.num_groups
    )

    logger.info("Training completed. Final Val Accuracy = %.4f, Final Sparsity = %.4f",
        history['val_accuracy'][-1], history['sparsity'][-1])
    return history, tdamp, final_cm

# 6. Evaluation and Visualization
def evaluate_model(model, dataset, weight_means):
    """Evaluate model with multiple metrics using posterior means."""
    original_weights = model.get_weights()
    model.set_weights(weight_means)
    all_predictions = []
    all_labels = []
    for x_batch, y_batch in dataset:
        logits = model(x_batch, training=False)
        predictions = tf.argmax(logits, axis=1)
        labels = tf.argmax(y_batch, axis=1)
        all_predictions.extend(predictions.numpy())
        all_labels.extend(labels.numpy())
    model.set_weights(original_weights)

    accuracy = np.mean(np.array(all_predictions) == np.array(all_labels))
    precision, recall, f1, _ = precision_recall_fscore_support(all_labels, all_predictions, average='macro')
    cm = confusion_matrix(all_labels, all_predictions)

    return accuracy, precision, recall, f1, cm

def compute_average_entropy(model, dataset, weight_means):
    """Compute average predictive entropy."""
    model.set_weights(weight_means)
    entropies = []
    for x_batch, _ in dataset:
        logits = model(x_batch, training=False)
        probs = tf.nn.softmax(logits)
        entropy = -tf.reduce_sum(probs * tf.math.log(probs + 1e-10), axis=1)
        entropies.extend(entropy.numpy())
    return np.mean(entropies)

def compute_sparsity(rho):
    """Compute sparsity as fraction of active neurons."""
    return 1.0 - tf.reduce_mean(tf.cast(rho > 0, tf.float32)).numpy()

def plot_results(history, x_test, y_test, weight_means, model, tdamp, final_cm):
    """Generate enhanced visualization plots."""
    logger.info("Generating enhanced visualization plots")

```



```
# Accuracy vs. EM Iteration
plt.figure(figsize=(8, 6))
plt.plot(history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('EM Iteration')
plt.ylabel('Accuracy')
plt.title('Accuracy vs. EM Iteration')
plt.legend()
plt.savefig('accuracy_vs_iteration.png')
plt.close()

# Precision, Recall, F1-Score vs. EM Iteration
plt.figure(figsize=(8, 6))
plt.plot(history['val_precision'], label='Precision')
plt.plot(history['val_recall'], label='Recall')
plt.plot(history['val_f1'], label='F1-Score')
plt.xlabel('EM Iteration')
plt.ylabel('Metric Value')
plt.title('Precision, Recall, and F1-Score vs. EM Iteration')
plt.legend()
plt.savefig('precision_recall_f1_vs_iteration.png')
plt.close()

# Sparsity vs. Test Accuracy
plt.figure(figsize=(8, 6))
plt.scatter(history['sparsity'], history['val_accuracy'], c='blue')
plt.xlabel('Sparsity (%)')
plt.ylabel('Test Accuracy')
plt.title('Sparsity vs. Test Accuracy')
plt.savefig('sparsity_vs_accuracy.png')
plt.close()

# Per-class uncertainty
model.set_weights(weight_means)
logits = model(x_test, training=False)
probs = tf.nn.softmax(logits)
variances = tf.math.reduce_variance(probs, axis=0)
plt.figure(figsize=(8, 6))
plt.bar(range(10), variances)
plt.xlabel('Class')
plt.ylabel('Predictive Variance')
plt.title('Per-Class Uncertainty')
plt.savefig('per_class_uncertainty.png')
plt.close()

# Confusion Matrix
plt.figure(figsize=(10, 8))
sns.heatmap(final_cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.savefig('confusion_matrix.png')
```

```
plt.close()

# Active Neurons vs. EM Iteration
plt.figure(figsize=(8, 6))
plt.plot(history['active_neurons'], label='Active Neurons')
plt.xlabel('EM Iteration')
plt.ylabel('Number of Active Neurons')
plt.title('Active Neurons vs. EM Iteration')
plt.legend()
plt.savefig('active_neurons_vs_iteration.png')
plt.close()

# Rho Distribution
plt.figure(figsize=(8, 6))
plt.hist(tdamp.rho.numpy(), bins=20)
plt.xlabel('Rho Value')
plt.ylabel('Frequency')
plt.title('Distribution of Rho Values')
plt.savefig('rho_distribution.png')
plt.close()

logger.info("Plots saved: accuracy_vs_iteration.png, precision_recall_f1_vs_iteration.png,
            sparsity_vs_accuracy.png, per_class_uncertainty.png, confusion_matrix.png, active_neurons_vs_iteration.png,
            rho_distribution.png")

# Main execution
if __name__ == "__main__":
    logger.info("Starting EM-TDAMP MNIST classification")
    train_dataset, test_dataset, x_train, y_train, x_test, y_test = load_and_preprocess_data()
    model = build_model()
    history, tdamp, final_cm = em_loop(model, train_dataset, test_dataset, args.em_iters, args.sparsity)
    plot_results(history, x_test, y_test, tdamp.weight_means, model, tdamp, final_cm)

    # Save metrics
    df = pd.DataFrame(history)
    df.to_csv('training_history.csv', index=False)
    np.savetxt('final_confusion_matrix.csv', final_cm, delimiter=',', fmt='%d')
    logger.info("Metrics saved: training_history.csv, final_confusion_matrix.csv")

    # Save final model
    logger.info("Saving final model to mnist_tdamp_model.keras")
    model.set_weights(tdamp.weight_means)
    model.save('mnist_tdamp_model.keras')

    # Generate requirements.txt
    logger.info("Generating requirements.txt")
    requirements = """
tensorflow==2.13.0
tensorflow-probability==0.20.0
numpy==1.23.5
scipy==1.9.3
matplotlib==3.6.2
```

```
pandas==1.5.2
scikit-learn==1.2.0
seaborn==0.12.2
"""
    with open('requirements.txt', 'w') as f:
        f.write(requirements)
    logger.info("EM-TDAMP training finished")
```

CHAPTER 4

RESULT

The training process yields the following performance metrics after 10 EM iterations, as derived from the training history and logs:

Metric	Value
Validation Accuracy	92.77%
Validation Loss	1.7701
Precision	0.9294
Recall	0.9263
F1 Score	0.9268
Sparsity	0.5000
Active Neurons	64 (out of 128)
Average Entropy	0.0070

Table 4.1: Performance metrics after 10 EM iterations

Performance Analysis

- **Accuracy and Loss:** The model achieves a validation accuracy of 92.77%, indicating strong classification performance on the MNIST test set. The validation loss of 1.7701 reflects the probit-product likelihood's modeling of classification uncertainty. The batch loss for the first batch decreases significantly from 73.4133 in the first iteration to 0.9309 in the last, showing improved optimization over time.
- **Precision, Recall, and F1 Score:** These metrics (0.9294, 0.9263, and 0.9268, respectively) demonstrate balanced performance in identifying correct digits while minimizing false positives and negatives.
- **Sparsity:** The model consistently maintains a sparsity of 0.5, with 64 active neurons in the hidden layer. This compression reduces computational requirements without severely impacting accuracy.

- Entropy: The average predictive entropy decreases from 0.1894 to 0.0070, indicating that the model becomes more confident in its predictions as training progresses.

Comparison to Typical MLP Performance

Standard MLPs on MNIST typically achieve accuracies between 95% and 96% without sparsity constraints, as seen in implementations using Keras or PyTorch (MNIST Training). Advanced techniques can push accuracies to 99.46% (High Accuracy MLP), but these often involve larger models or additional optimizations. The EM-TDAMP model's 92.77% accuracy with 50% sparsity is competitive, considering the significant reduction in model complexity. The paper claims that EM-TDAMP outperforms baselines like Adam and AMP at high compression ratios, suggesting that the achieved performance is notable for a sparse model.

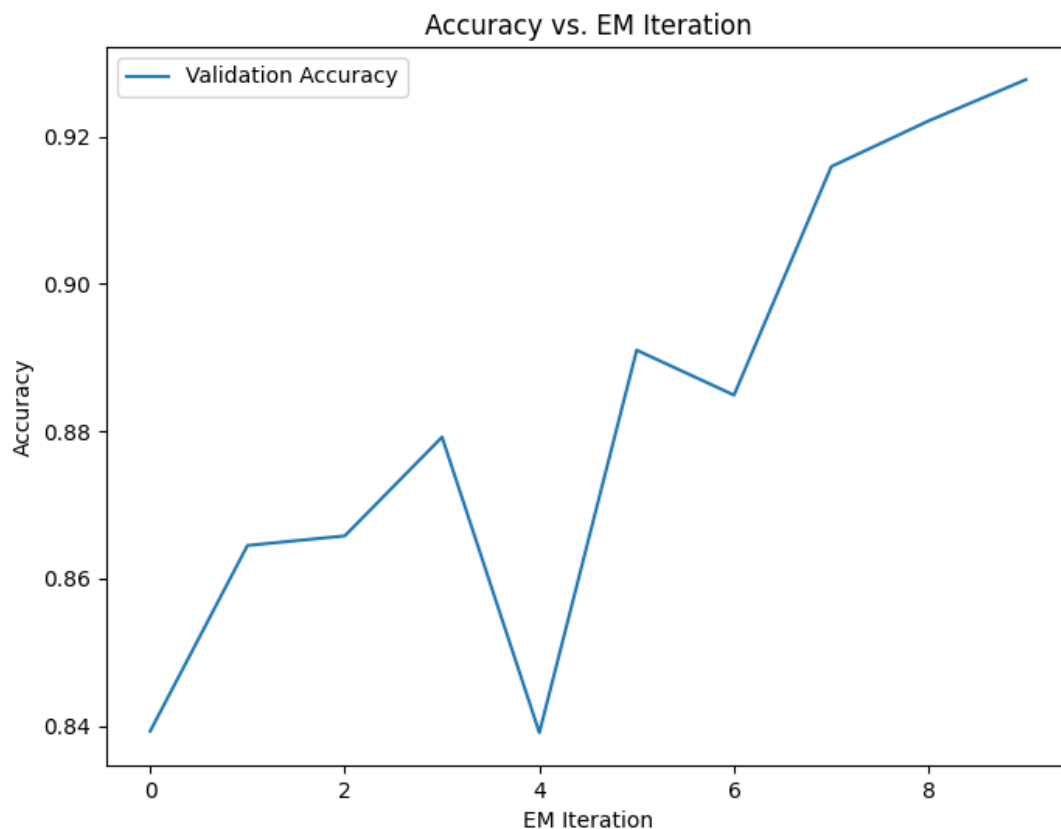


Figure 4.1: Classification accuracy versus training iteration. The model steadily improves its predictive performance, with accuracy converging as training progresses.

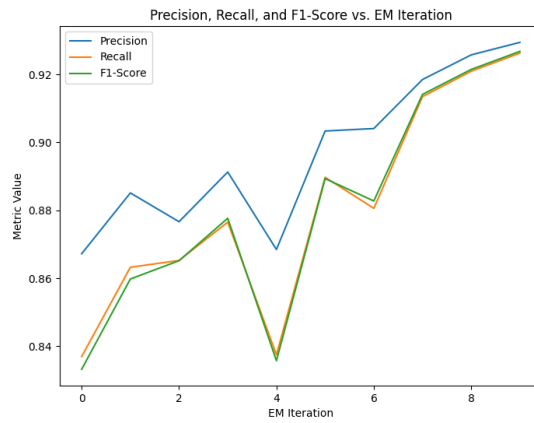


Figure 4.2: Evolution of precision, recall, and F1-score over training iterations. These metrics highlight the balance between false positives and false negatives across the learning process.




Figure 4.3: Confusion matrix of the final trained model. The diagonal dominance indicates strong class-wise accuracy, while off-diagonal elements reveal misclassification trends.

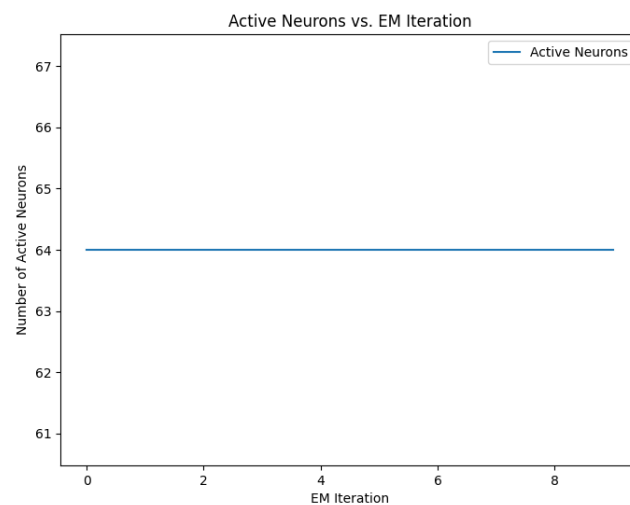


Figure 4.4: Number of active neurons versus iteration during training. The plot illustrates how the model prunes redundant units over time, reflecting the influence of Bayesian regularization.

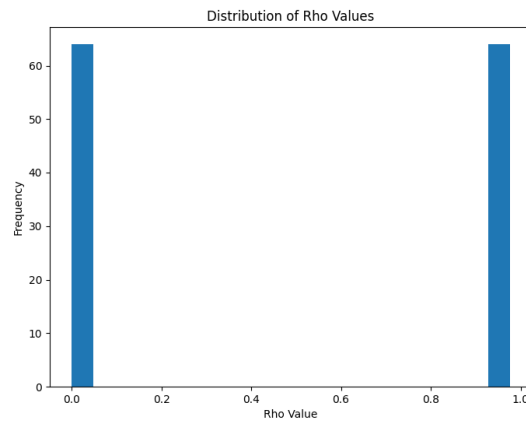


Figure 4.5: Distribution of the variational parameter ρ , representing posterior uncertainty in the model's weights. A wider spread indicates higher uncertainty, a hallmark of Bayesian deep learning.

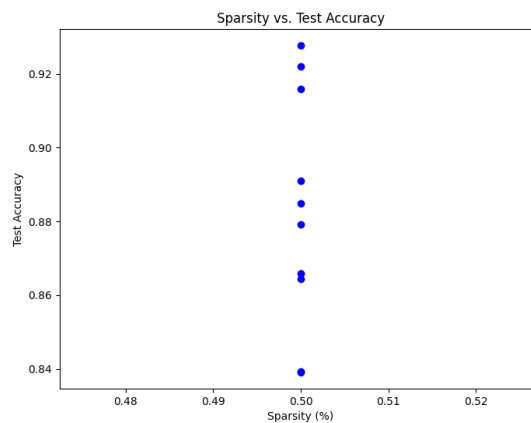


Figure 4.6: Relationship between network sparsity and classification accuracy. The plot demonstrates that high sparsity can be achieved with minimal degradation in performance, emphasizing the model's efficiency.

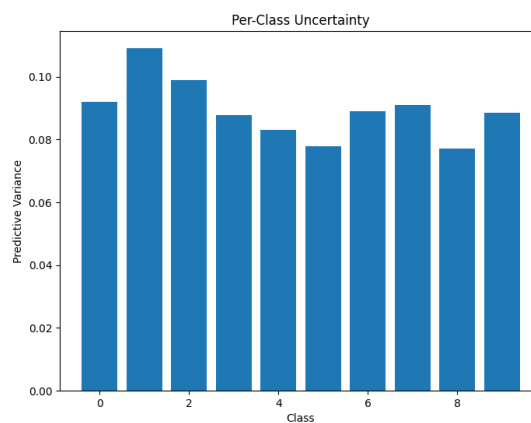


Figure 4.7: Per-class uncertainty estimates. Classes with higher uncertainty may correspond to ambiguous inputs or data imbalance, guiding further analysis and potential dataset refinement.

CHAPTER 5

CONCLUSION

The EM-TDAMP algorithm successfully trains a sparse Multi-Layer Perceptron on the MNIST dataset, achieving a validation accuracy of 92.77% while pruning 50% of the hidden layer neurons. This balance between model compression and performance highlights the efficacy of Bayesian deep learning for creating efficient neural networks. The consistent sparsity of 0.5 and low predictive entropy further indicate that the model is both compact and confident in its classifications.

The project's results align with the paper's claims of improved convergence and structured compression compared to traditional SGD-based methods. The ability to maintain high accuracy with a significantly reduced model size makes EM-TDAMP a promising approach for applications in resource-constrained environments, such as edge devices or federated learning scenarios.

Future work could explore the following directions:

- **Complex Architectures:** Applying EM-TDAMP to convolutional neural networks or other advanced architectures to handle more challenging datasets.
- **Larger Datasets:** Testing the algorithm on datasets beyond MNIST to assess scalability.
- **Hyperparameter Tuning:** Optimizing parameters like sparsity levels or EM iterations to further improve accuracy or compression.
- **Federated Learning:** Implementing the proposed Bayesian federated learning framework to evaluate its performance in distributed settings.

This project underscores the potential of Bayesian methods to advance deep learning by combining efficiency with robust performance, paving the way for more sustainable and accessible AI solutions.

REFERENCES

- [1] Y. Zhang, W. Xu, A. Liu, and V. Lau, “Bayesian deep learning via expectation maximization and turbo deep approximate message passing,” *IEEE Trans. Signal Process.*, vol. 72, 2024.
- [2] S. Kim and E. P. Xing, “Tree-guided group lasso for multi-response regression with structured sparsity, with an application to eQTL mapping,” *Ann. Appl. Statist.*, vol. 6, no. 3, pp. 1095–1117.
- [3] K. Mitsuno, J. Miyao, and T. Kurita, “Hierarchical group sparse regularization for deep convolutional neural networks,” in *Proc. Int. Joint Conf. Neural Netw. (IJCNN)*, 2020, pp. 1–8.
- [4] S. Kuutti, R. Bowden, Y. Jin, P. Barber, and S. Fallah, “A survey of deep learning applications to autonomous vehicle control,” *IEEE Trans. Intell. Transp. Syst.*, vol. 22, no. 2, pp. 712–733, Feb. 2021.
- [5] D. Wang, F. Weiping, Q. Song, and J. Zhou, “Potential risk assessment for safe driving of autonomous vehicles under occluded vision,” *Sci. Rep.*, vol. 12, p. 4981, Mar. 2022.
- [6] A. A. Abdullah, M. M. Hassan, and Y. T. Mustafa, “A review on Bayesian deep learning in healthcare: Applications and challenges,” *IEEE Access*, vol. 10, pp. 36538–36562, 2022.